



# Managing Real-Time Media Flows through a Flow State Exchange

Safiqul Islam, Michael Welzl, David Hayes, Stein Gjessing

*Networks and Distributed Systems Group*

*Department of Informatics*

*University of Oslo*

# Outline

- Context
- Problem Statement
- Example Architecture
- Flow State Exchange (FSE)
- Results
- Summary

# Context/Background

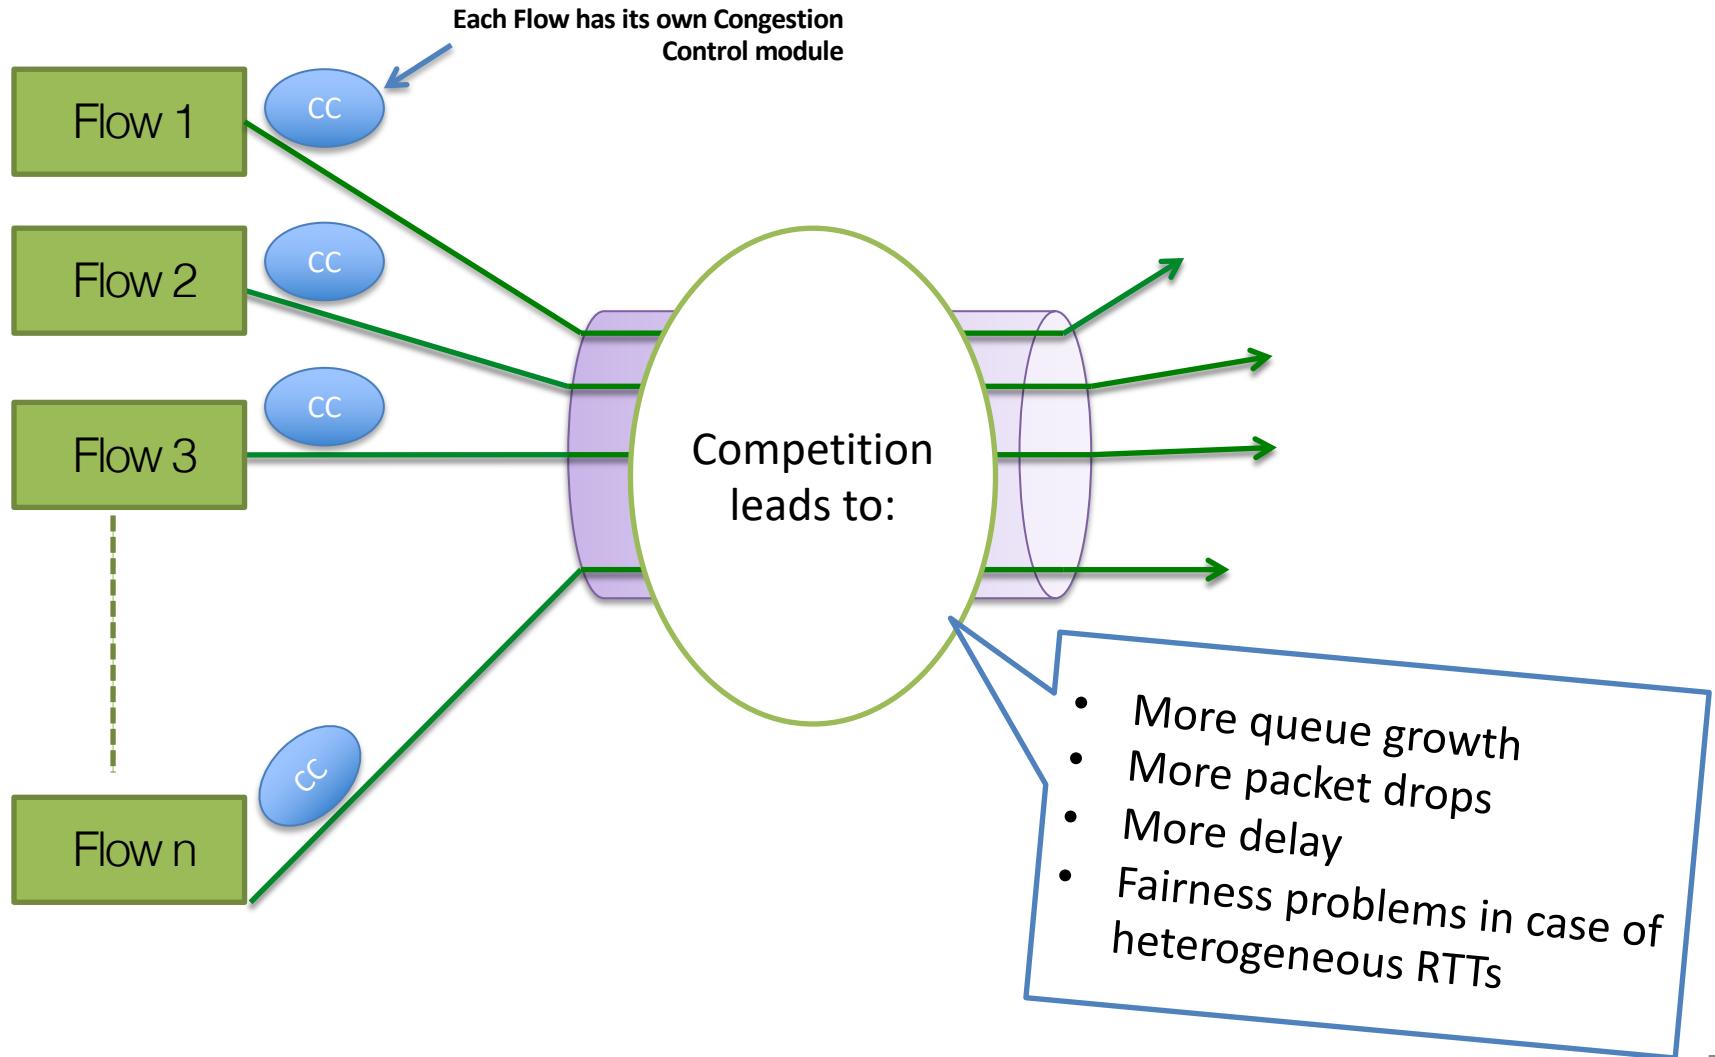
- WebRTC/RTCWeb enables real-time communication between web-browsers.



# Context/Background

- Congestion control for real-time media is delay sensitive
  - However, Except TFRC, there are no standardized protocols until now
- IETF RMCAT WG is formed to develop standards for RTP based Media
- Three mechanisms are WG items
  - Cisco's Network Assisted (NADA)
  - Google's Google Congestion Control
  - Ericsson's ScreeAm

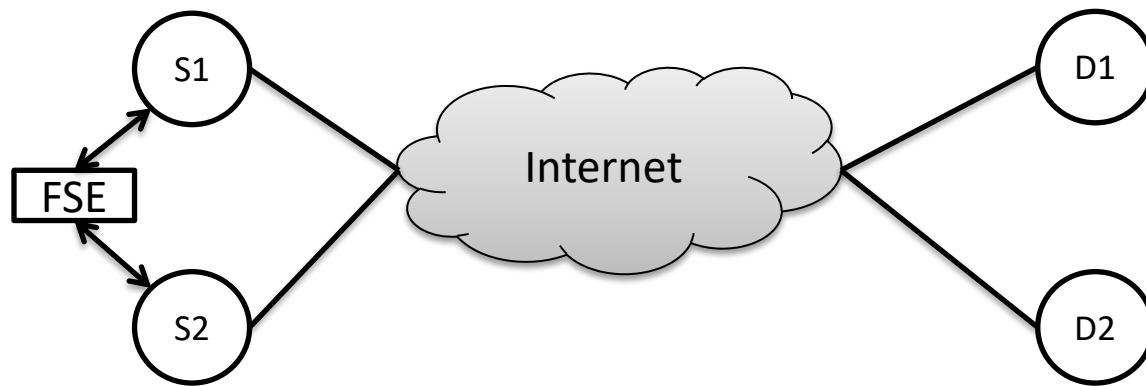
# Problem Statement



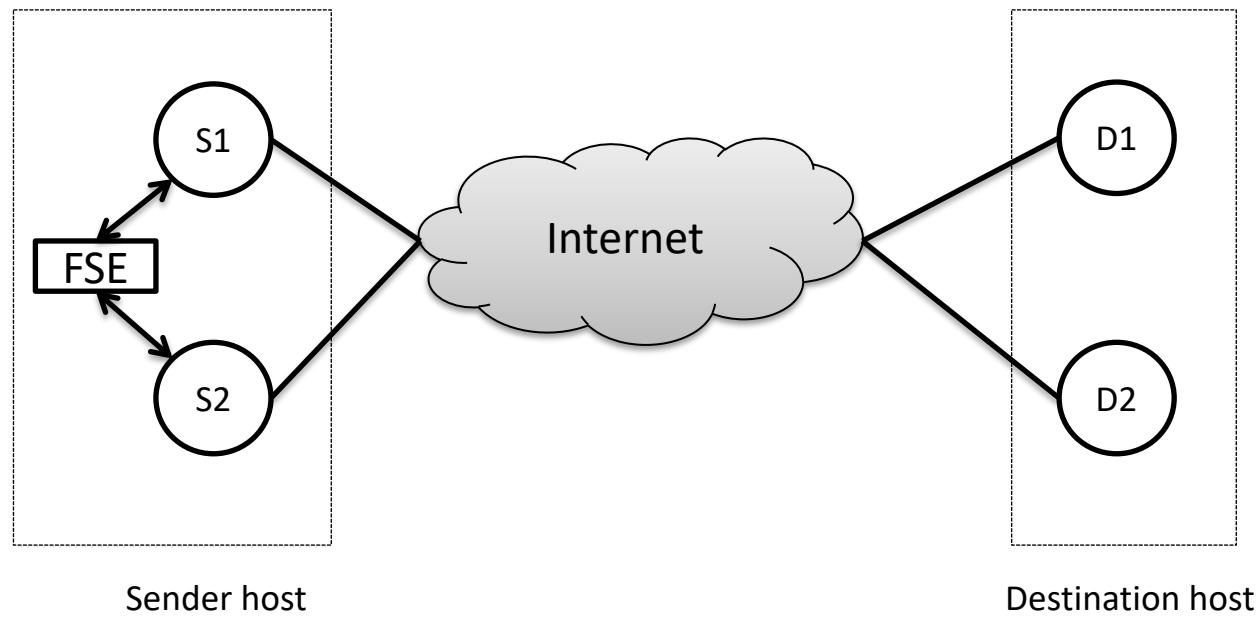
# Problem Statement

- This can be solved with a single congestion controller, but it's too complicated
  - Congestion Manager
- Ensemble-TCP, EFCM
  - Only for TCP, and requires Extensive changes

# Example Architecture

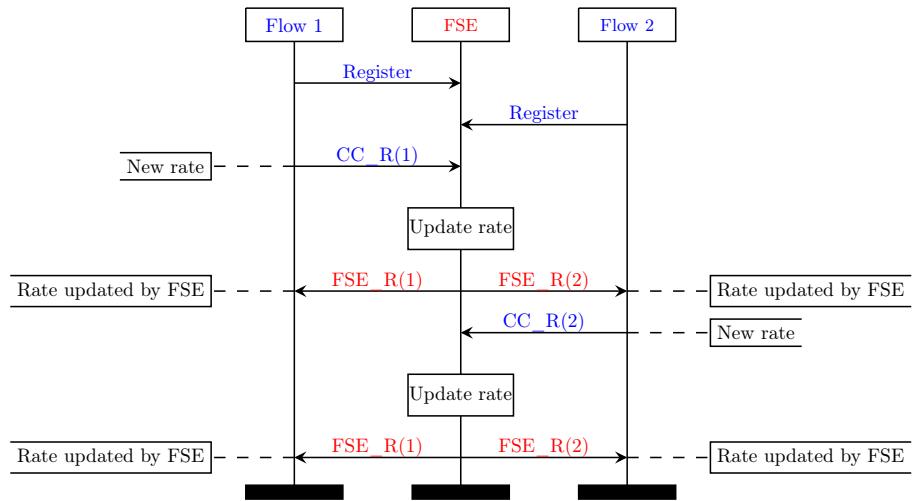


# Example Architecture

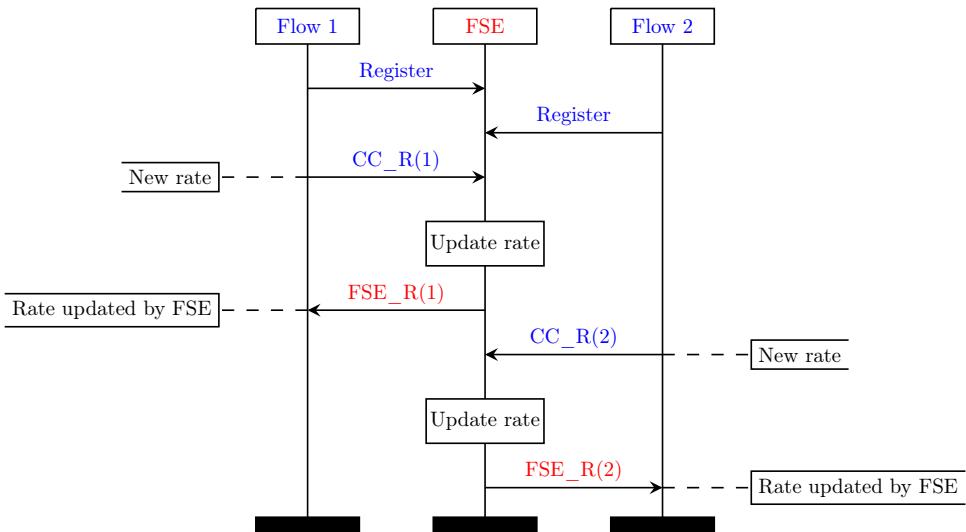


# Versions?

msc Active FSE

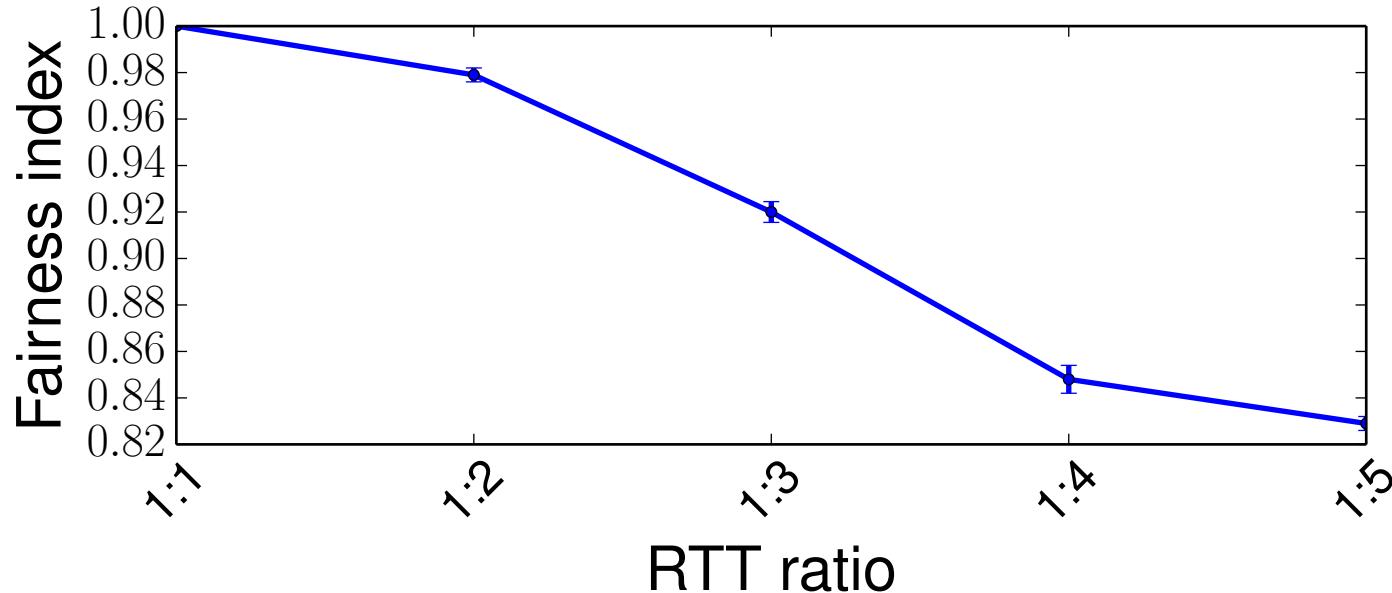


msc Passive FSE



# Why Passive version doesn't always work?

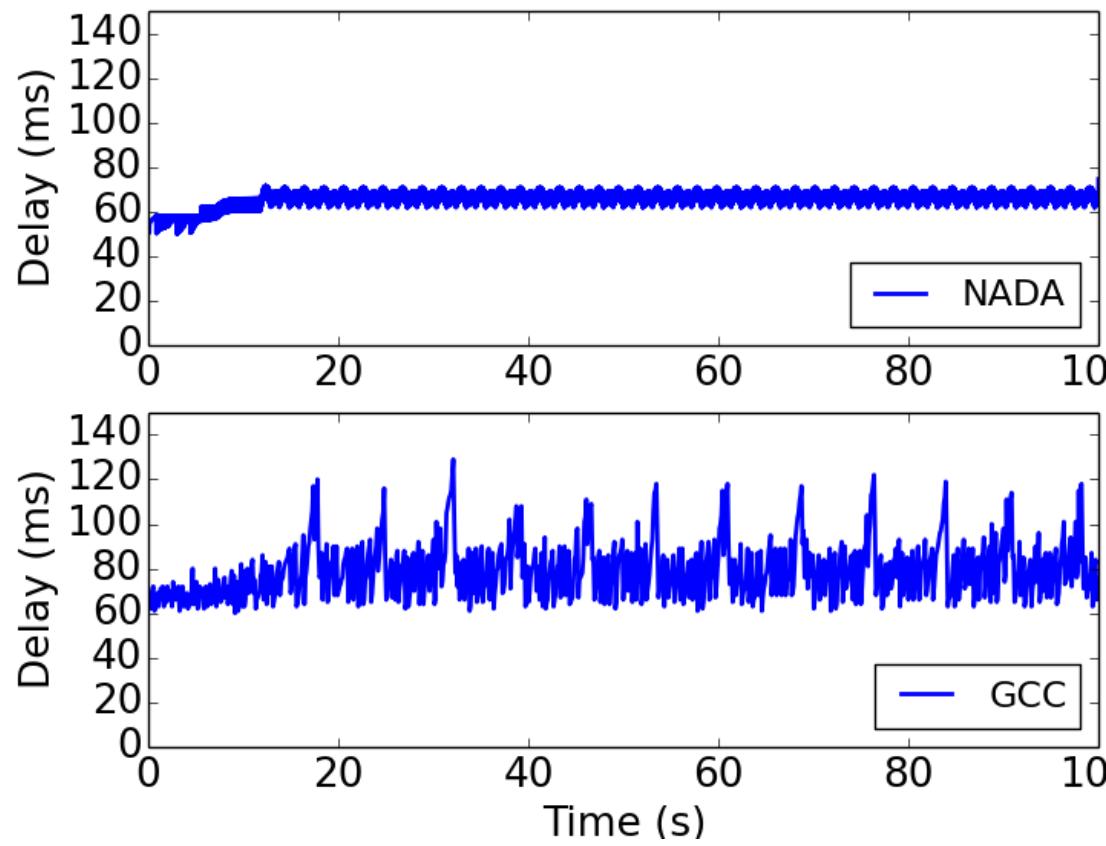
- ...



# Adaptation of the algorithm

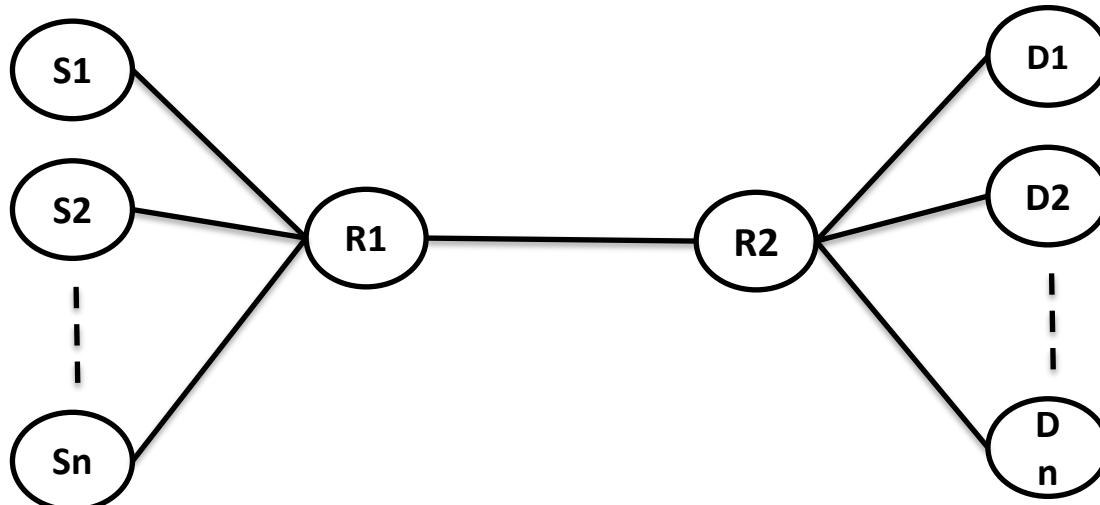
- ...

# Adaptation of the algorithm



# Simulation Testbed

- Implemented the mechanism in both Chromium Browser and ns-2
- Actual tests and media characteristics are hard to predict
  - IETF RMCAT group prepared a number of test-cases
    - Dumbbell topology
    - With a bottleneck queue of 300ms
    - We show two pertinent test cases: Round-trip time fairness and media pause and resume
  - Prioritization test
  - Delayed feedback test



# Prioritization Test

- Controllable fairness with prioritization is one of the key requirements of WebRTC

Have you personally experienced that the network traffic of applications on your computer have influenced each other?

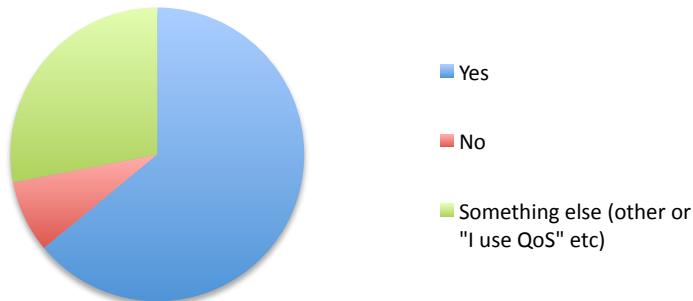


Yes, and I found it annoying

Yes, but I didn't care

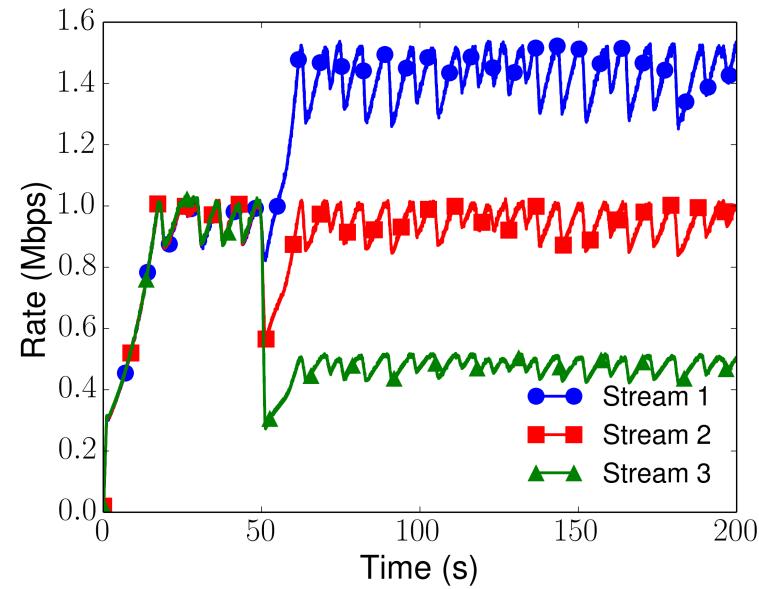
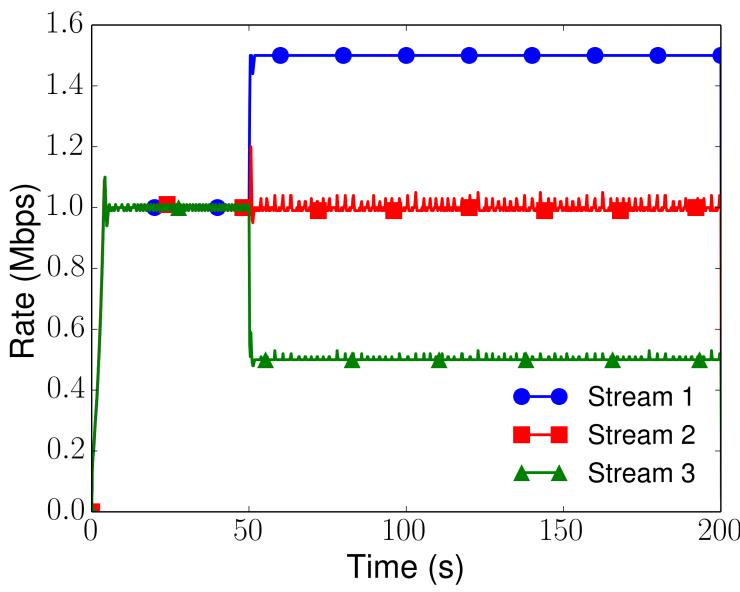
No: this never happened - or if it did, I didn't notice

If there was an easy-to-use tool that would let me prioritize how my applications access the network, I'd use it



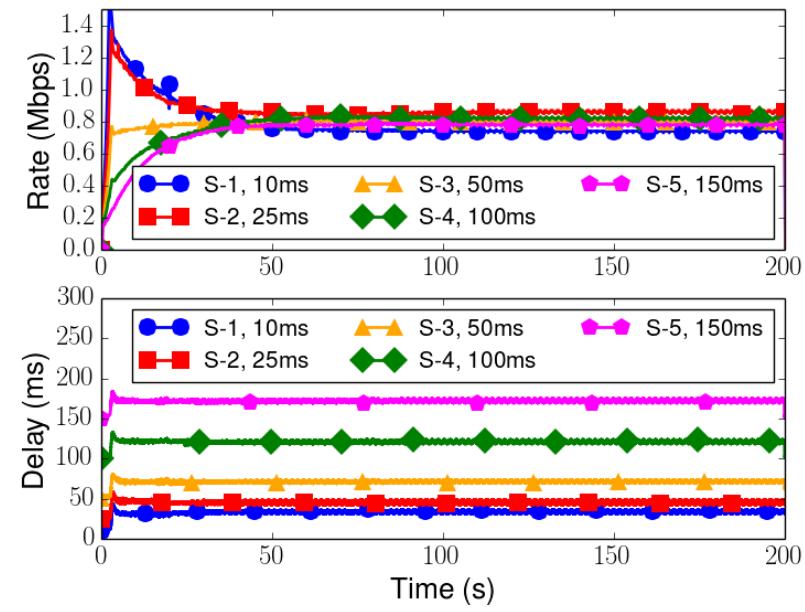
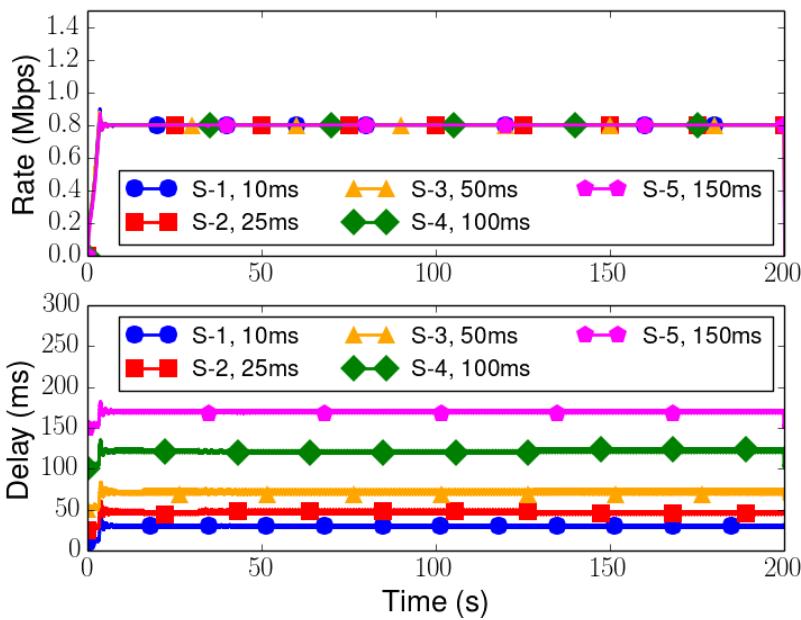
# Prioritization Test

- The FSE can assign rates based on the priorities assigned from the uses
- *One way base delay, 50ms*



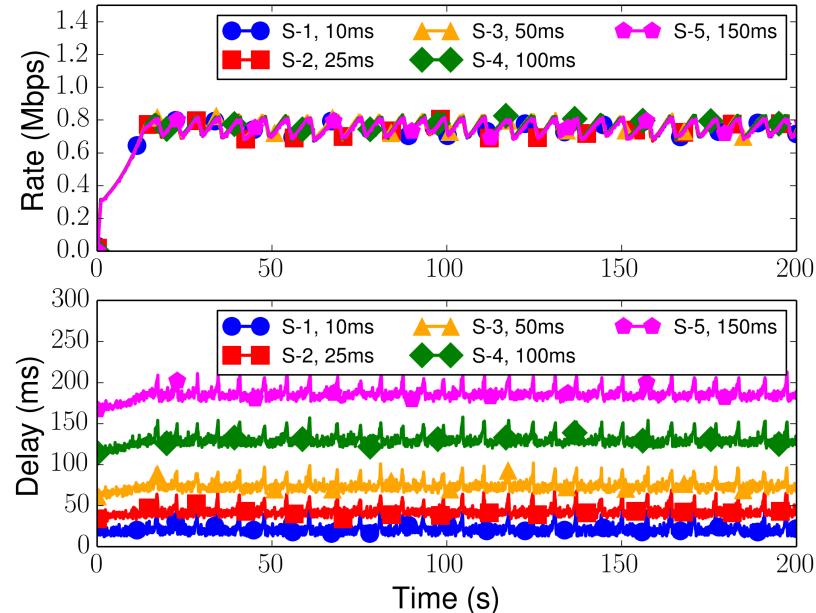
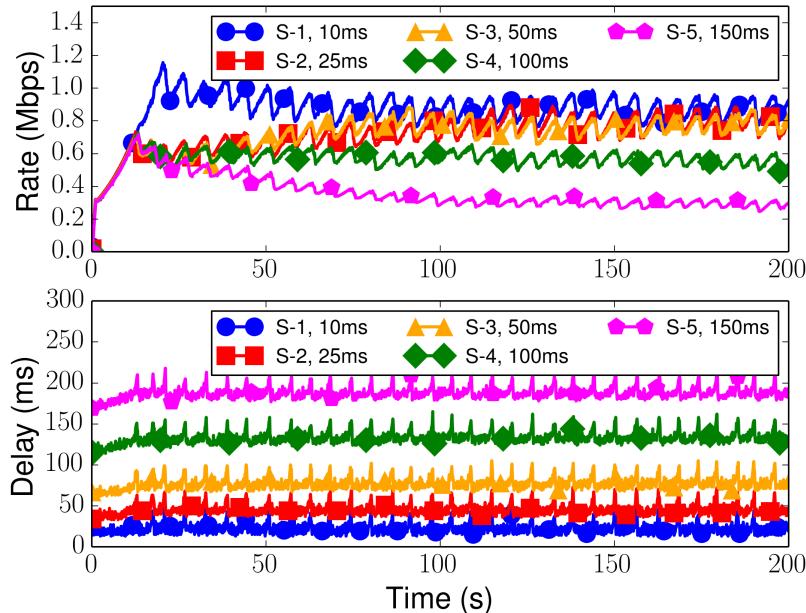
# Round-trip time fairness, NADA flows

- One-way delays of s1, s2, s3, s4 and s5 are 10ms, 25ms, 50ms, 100ms, and 150ms respectively, and bottleneck capacity, 4Mbps



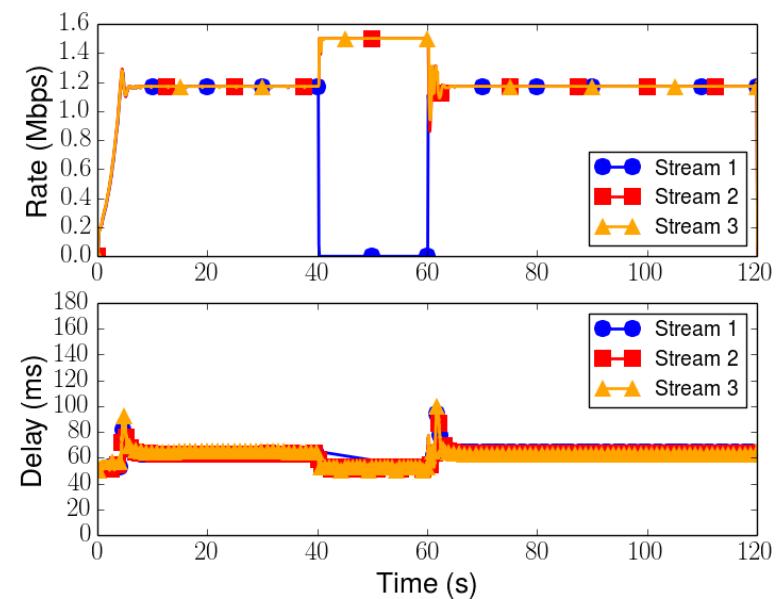
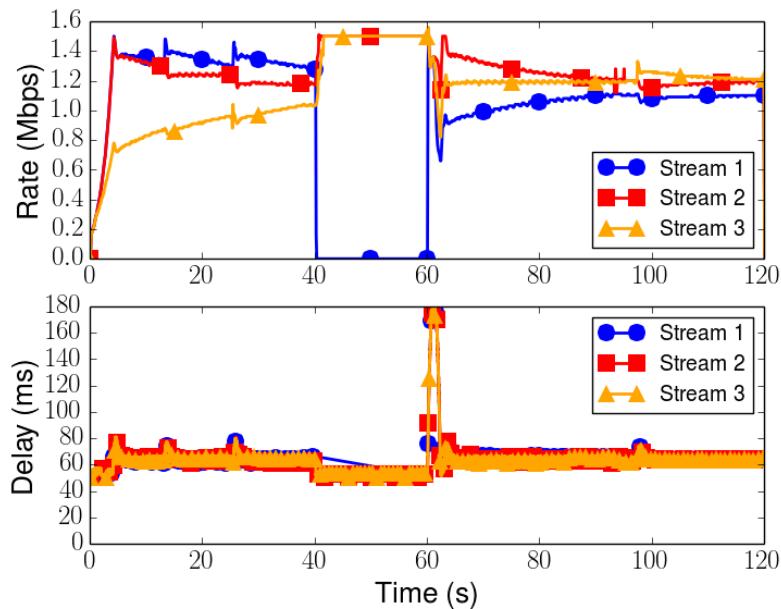
# Round-trip time fairness – GCC flows

- One-way delays of s1, s2, s3, s4 and s5 are 10ms, 25ms, 50ms, 100ms, and 150ms respectively, and bottleneck capacity, 4Mbps



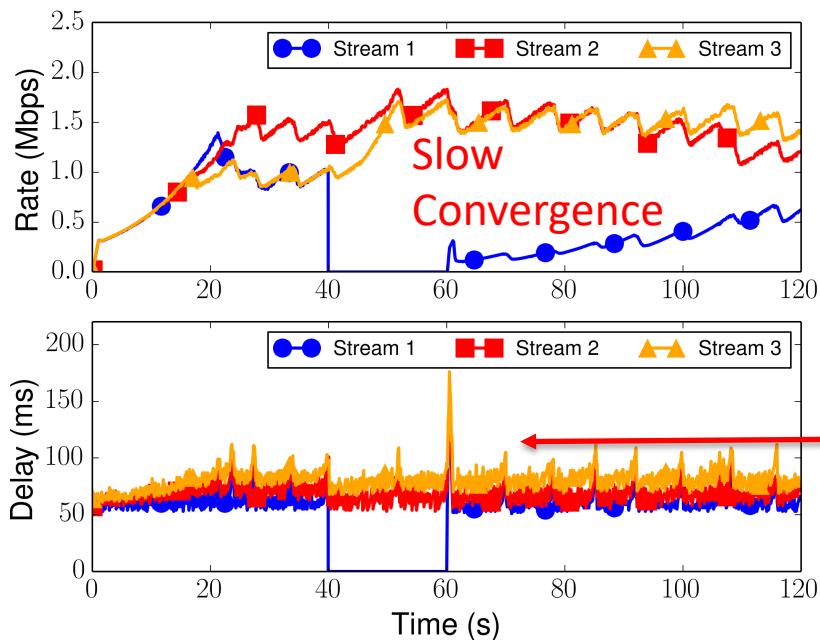
# Media Pause and Resume (NADA)

- One way delay 50ms and Path capacity 3.5 Mbps

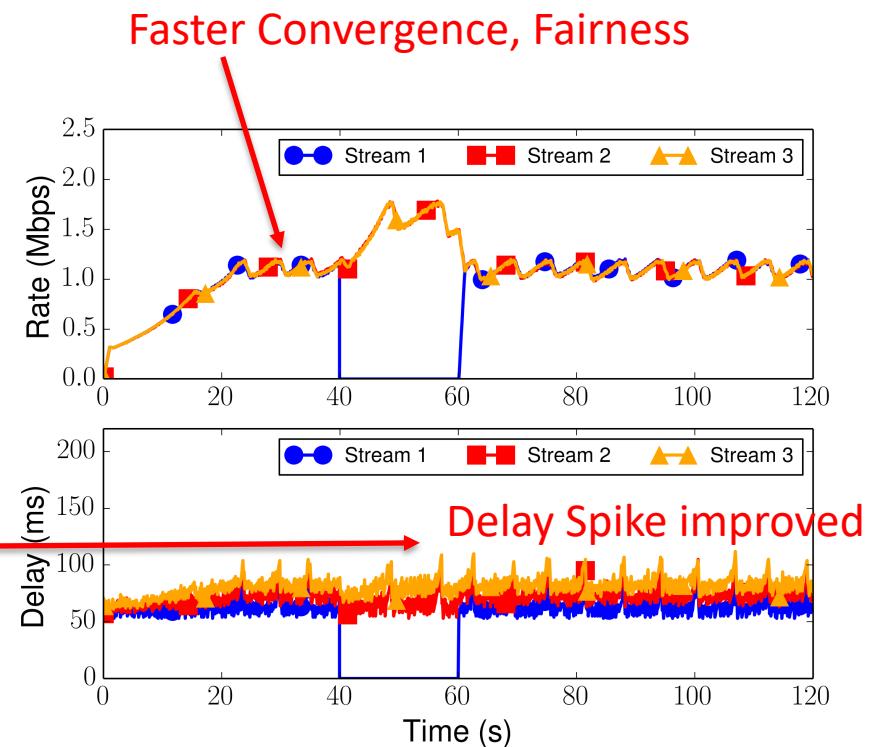


# Media Pause and Resume (GCC)

- One way delay 50ms and Path capacity 3.5 Mbps



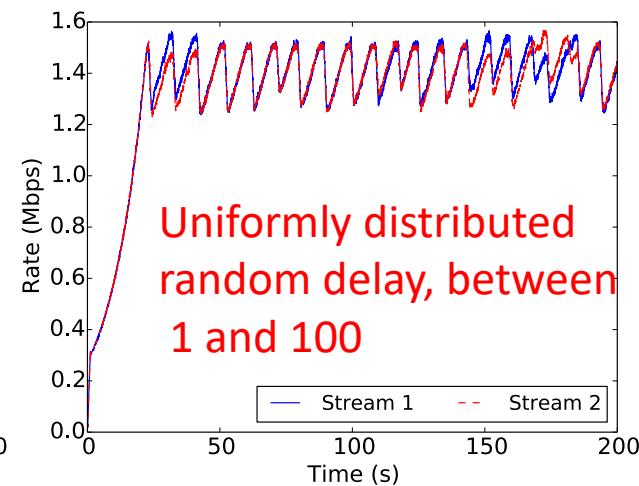
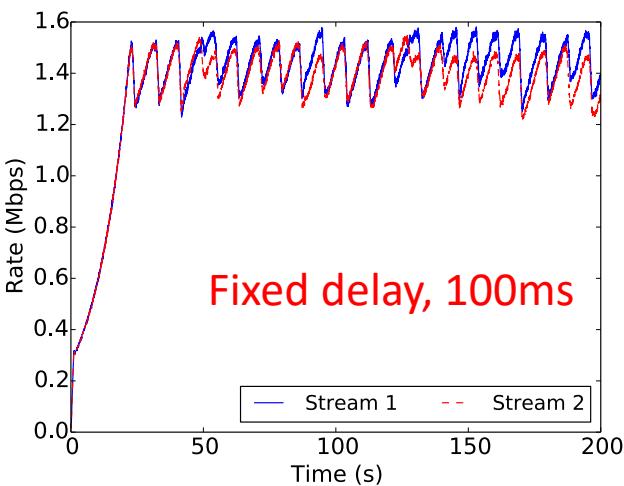
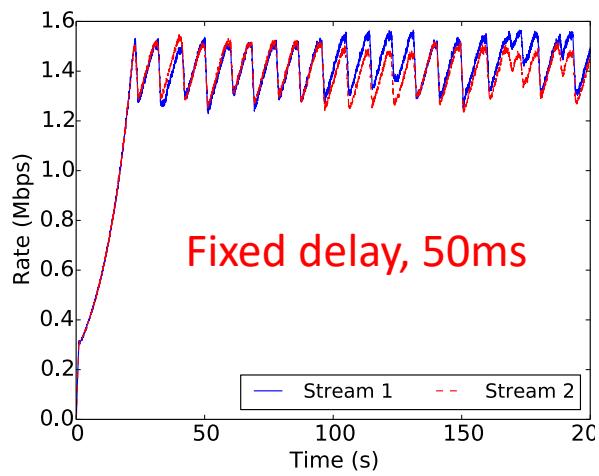
Without FSE



With FSE

# Delayed feedback test

- Delay between stream 1 and the FSE is delayed



# Q&A